

Fine-mapped single-cell eQTL evaluation of AlphaGenome fine-tuning

1 Overview

We evaluate whether sequence-derived mutation effects distinguish fine-mapped single-cell expression quantitative trait loci (eQTLs). A positive is a variant with posterior inclusion probability (PIP) greater than 0.75 and a negative is a variant with PIP less than 0.01. The evaluation uses human fine-mapping panels and a 1 Mb sequence window centered on the transcription start site (TSS) of the associated gene. RNA-seq denotes RNA sequencing and CAGE denotes Cap Analysis of Gene Expression. Average precision (AP) is reported as area under the precision-recall curve (AUPR); AUROC denotes area under the receiver operating characteristic curve.

The primary comparison is a trained head-only model versus the two best available joint fine-tuned models, a low-rank adaptation (LoRA) model and a LoRA model with convolutional low-rank adapters (LoCon). The benchmark retains up to 1,000 variants in each PIP class per fine-mapping panel with a fixed random seed. The raw and retained counts are recorded with the benchmark output.

2 Cell-type matching

The 35 fine-mapping panels use the prefixes Ast (astrocyte), End (endothelial), Ext (excitatory), IN (inhibitory), MG (microglia), OD (oligodendrocyte), and OPC (oligodendrocyte precursor cell), with numbered panels denoting subpanels. The trained head-only, LoRA, and LoRA+LoCon models use the same target track definitions. The matching used for these models is shown below. A separate native AlphaGenome analysis evaluates all valid pretrained RNA-seq and CAGE tracks, without adding a newly initialized head.

Table 1: Fine-mapping panel to RNA-track matching for the trained head-only, LoRA, and LoRA+LoCon models. A numbered panel inherits the mapping for its prefix.

Panel prefix	Trained head-only	LoRA and LoCon
Ast*	ASC	ASC
End*	Endo	Endo
Ext*	Mean of L2/3 IT, L4/5 IT, L5/6 NP, L5 ET, L5 IT, L6 CT, L6 IT CAR3, and L6b	Same
IN*	Mean of LAMP5, PVALB, SNCG, SST, and VIP	Same
MG*	MGC	MGC
OD*	ODC	ODC
OPC*	OPC	OPC

The panel labels do not exactly match the available pretrained output tracks. The native model therefore receives a separate track-level analysis; it is not represented by the trained-head comparison above.

3 Mutation effects

For variant i , let $\mathbf{y}_{i,r} \in \mathbb{R}^{B \times C}$ and $\mathbf{y}_{i,a} \in \mathbb{R}^{B \times C}$ be the reference and alternate RNA predictions at the model’s 128 base-pair (bp) resolution, where B is the number of output bins and $C = 19$ is the number of RNA tracks. Let $\mathcal{C}_g$ be the track indices for the matched cell group. The TSS score is

$$s_i^{\text{TSS}} = \frac{1}{|\mathcal{C}_g|} \sum_{c \in \mathcal{C}_g} |y_{i,a,b_{\text{TSS}},c} - y_{i,r,b_{\text{TSS}},c}|, \quad (1)$$

where b_{TSS} is the 128 bp bin containing the TSS. The gene score is the corresponding mean absolute alternate-reference difference after averaging each track over all 128 bp bins overlapping the annotated gene span,

$$s_i^{\text{gene}} = \frac{1}{|\mathcal{C}_g|} \sum_{c \in \mathcal{C}_g} \left| \frac{1}{|\mathcal{B}_i|} \sum_{b \in \mathcal{B}_i} y_{i,a,b,c} - \frac{1}{|\mathcal{B}_i|} \sum_{b \in \mathcal{B}_i} y_{i,r,b,c} \right|, \quad (2)$$

where $\mathcal{B}_i$ is the set of bins overlapping the gene span. Both scores are nonnegative because fine-mapping identifies variant probability, not the direction of expression change. Performance is measured by AUROC and average precision, overall and within bins of $|\text{variant} - \text{TSS}|$.

4 Results

Each trained model scored the same 54,193 variants after reference-allele validation. LoRA has the highest TSS-bin AUROC and AP; the trained head-only checkpoint has the highest gene-span AUROC and AP. Overall differences are small, and the distance-stratified curves show similar trends.

Table 2: Overall eQTL discrimination across scored variants. AP is average precision.

Model	TSS AUROC	TSS AP	Gene AUROC	Gene AP
Head-only	0.5109	0.3817	0.5394	0.4269
LoRA	0.5155	0.3862	0.5329	0.4156
LoRA+LoCon	0.5091	0.3831	0.5233	0.4131

Performance by distance from the TSS is shown below. Per-bin AUROC, AP, and class counts are available in the machine-readable summary output.

4.1 Native track specificity

The native AlphaGenome evaluation stores signed alternate-minus-reference effects for 54,729 sampled variants and 1,213 valid pretrained channels, 667 RNA-seq and 546 CAGE. Predictions are reduced on device to a 128 bp TSS bin and the mean over the annotated gene span. Of the variants, 54,193 have finite effects and 536 were skipped because the reference allele did not match the reference genome. Absolute effects are scored separately for each biosample against the positive and negative variants in each panel group; strand and channel tracks are grouped by source, biosample, and assay. Ranks are among the 371 RNA-seq or 273 CAGE biosample groups. The full rank distribution is shown in Figure 3.

Table 3 lists the highest-AP native track across all 644 biosample/assay groups for each cell type, selecting the winner separately for the TSS-bin and gene-span summaries. These absolute winners




Figure 1: AUROC by distance from the TSS for TSS-bin and gene-span mutation effects. The dashed line is chance performance.

need not be biologically matched tracks; for example, testis RNA-seq is the top track for several brain-cell panels. The heatmap below instead focuses on annotated direct or plausible cell-type candidates.

Table 3: Highest average-precision native RNA-seq or CAGE track for each eQTL cell type. Winners are selected independently by summary from all RNA-seq and CAGE biosample groups.

eQTL cell type	Best TSS-bin track	AP	Best gene-span track	AP
Ast	Dorsolateral prefrontal cortex (ENCODE, total RNA-seq)	0.400	Testis (ENCODE, total RNA-seq)	0.459
End	Tibial nerve (ENCODE, total RNA-seq)	0.259	Testis (ENCODE, total RNA-seq)	0.294
Ext	Testis (ENCODE, total RNA-seq)	0.476	Testis (ENCODE, total RNA-seq)	0.554
IN	Testis (GTEx, polyA RNA-seq)	0.405	Testis (ENCODE, total RNA-seq)	0.486
MG	Spleen (ENCODE, total RNA-seq)	0.345	Spleen (ENCODE, total RNA-seq)	0.418
OD	Dorsolateral prefrontal cortex (ENCODE, total RNA-seq)	0.567	Dorsolateral prefrontal cortex (ENCODE, total RNA-seq)	0.607
OPC	Dorsolateral prefrontal cortex (ENCODE, total RNA-seq)	0.452	Parietal lobe (ENCODE, total RNA-seq)	0.522

Monocyte tracks are the clearest microglia proxy. On MG panels, generic monocyte CAGE ranks first of 273 biosample groups for TSS AP (0.320) and second for gene-span AP (0.308), versus unmatched-track medians of 0.261 and 0.242. The CD14-positive, CD16-negative classical monocyte CAGE track ranks seventh for TSS AP (0.306) and fourth for gene-span AP (0.295). CD14-positive monocyte RNA-seq ranks 27th of 371 for TSS AP (0.307) and 47th for gene-span AP (0.332). This supports CD14-positive monocytes as a useful prediction proxy for these eQTL panels, without equating monocytes with microglia biologically.

There is no mature oligodendrocyte-labeled native RNA-seq or CAGE channel. The single oligodendrocyte-precursor CAGE channel ranks 226th of 273 on OPC TSS AP (0.354) and 246th




Figure 2: AUPR by distance from the TSS for TSS-bin and gene-span mutation effects. Black diamonds show the random-ranking reference, equal to the positive variant fraction in each distance bin.

on gene-span AP (0.336), below the unmatched-track medians (0.362 and 0.345). Bulk brain tracks perform better: dorsolateral prefrontal cortex RNA-seq ranks first of 371 for OD panels on both summaries (AP 0.567 and 0.607), and brain-region tracks also rank highly for OPC panels. For OPC panels, dorsolateral prefrontal cortex RNA-seq ranks first for TSS AP and third for gene-span AP. These are broad neural-tissue proxies, not evidence of oligodendrocyte-specific prediction. Astrocyte tracks are weakly ranked; endothelial tracks are moderately enriched over unmatched tracks. The available glutamatergic and generic neuronal tracks provide mixed results, and there is no inhibitory-neuron-specific native channel.

Figure 4 shows average precision (AP) for candidate tracks across all seven eQTL cell types. Candidate rows are the union of tracks ranking in the top five annotated candidates for at least one cell type and one summary; AP is then shown for every cell type. Tracks are grouped by output type, biosample, source, and assay. Separate panels retain the TSS-bin and gene-span summaries. Because positive-class prevalence differs between cell-type panels, AP is most directly compared between tracks within the same column.




Figure 3: Average-precision rank of native tracks annotated as direct cell labels or plausible proxies within each fine-mapping panel. Rank 1 is best; RNA-seq and CAGE biosample groups are ranked separately. Bulk-brain ranks indicate tissue relevance and should not be interpreted as cell specificity.




Figure 4: Native AlphaGenome candidate-track average precision by eQTL cell type. Rows are the union of tracks ranked in the top five annotated candidates for at least one cell type and summary; columns show AP for each cell type.

These track ranks are descriptive rather than a formal statistical test of cell-type specificity. Native track availability and assay/source coverage differ across cell classes, and most fine-mapping labels have no exact native counterpart.

5 Implementation details

Fine-mapping tables are streamed from compressed PIP files. Variant alleles are checked against the reference genome before constructing the alternate one-hot sequence. The same reference and alternate sequences are evaluated for every model. The head-only checkpoint was trained on the target tracks without LoRA or LoCon adapters; the fine-tuned models retain trained heads and adapters. All model prediction calls use the custom forward path so the fine-tuned heads, rather than the delegated pretrained variant scorer, are evaluated. Native AlphaGenome variant effect predictions (VEPs) use the original model’s valid RNA-seq and CAGE tracks, with predictions reduced on device to the TSS bin and gene span before signed alternate-minus-reference effects are stored. The full matrices and channel/variant metadata accompany the analysis outputs. Candidate tracks were ranked against all biosample groups separately for each panel and summary; the plotted annotations distinguish direct labels, neuronal/monocyte proxies, oligodendrocyte-precursor tracks, and bulk brain tissue.